

Deep Learning

Loo Tiang Kuan, Leonard

A0098795B

supervised by

Associate Professor Ji Hui

co-supervised by

Assistant Professor Alvina Goh Siew Wee

May 5, 2017

Today's Agenda

- 1 Denoising
- 2 Autoencoders
- 3 Generative Adversarial Networks
- 4 Improving Discriminator
- 5 Summary & Interesting Applications

Outline

1 Denoising

- Beyond A Gaussian Denoiser
 - Blind Denoising

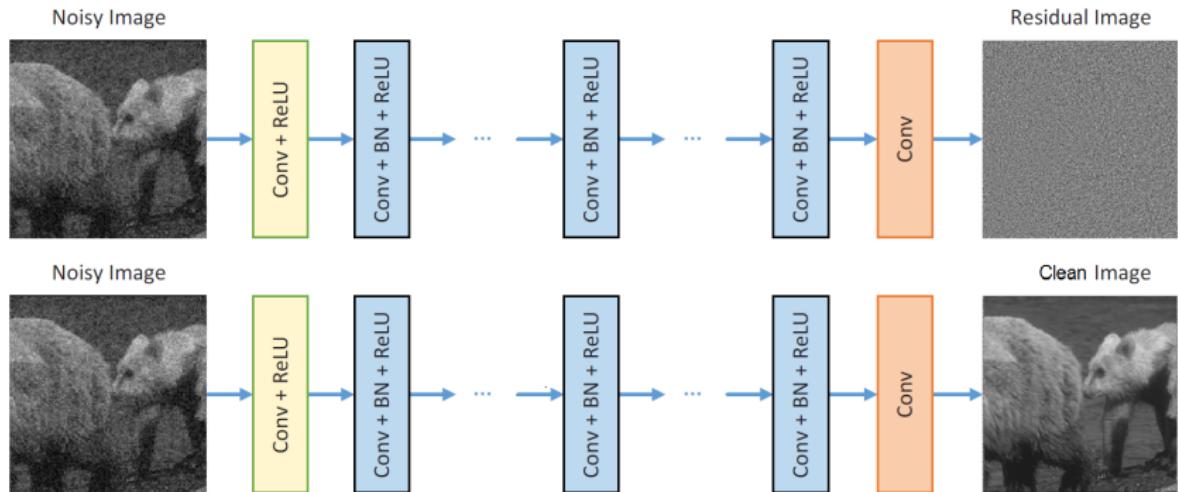
2 Autoencoders

3 Generative Adversarial Networks

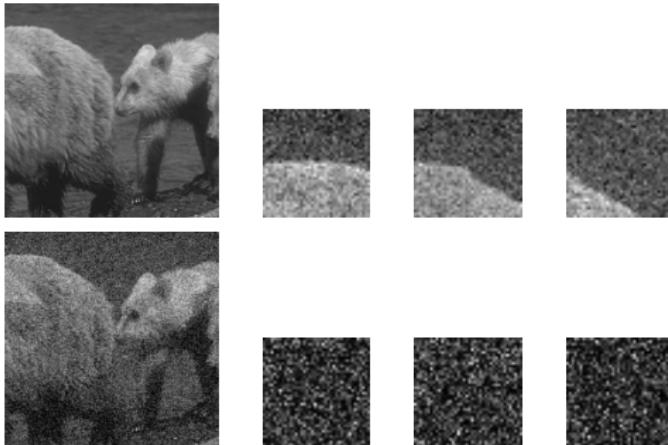
4 Improving Discriminator

5 Summary & Interesting Applications

Architecture



Training



- Motivation from ResNet (macro scale only)
- Can feed image of any size ('same' padding throughout).
- Heavy data augmentation

Experiments (Trained only with sigma=10)

sigma = 5

sigma = 10

sigma = 15

Top Row: Noisy

Bottom Row: Denoised

Outline

1 Denoising

- Beyond A Gaussian Denoiser
 - Blind Denoising

2 Autoencoders

3 Generative Adversarial Networks

4 Improving Discriminator

5 Summary & Interesting Applications

Blind Denoising



- Train with varying variance \implies blind denoising on range of the variances trained. $\sigma \in [5, 50]$

Sigma=10



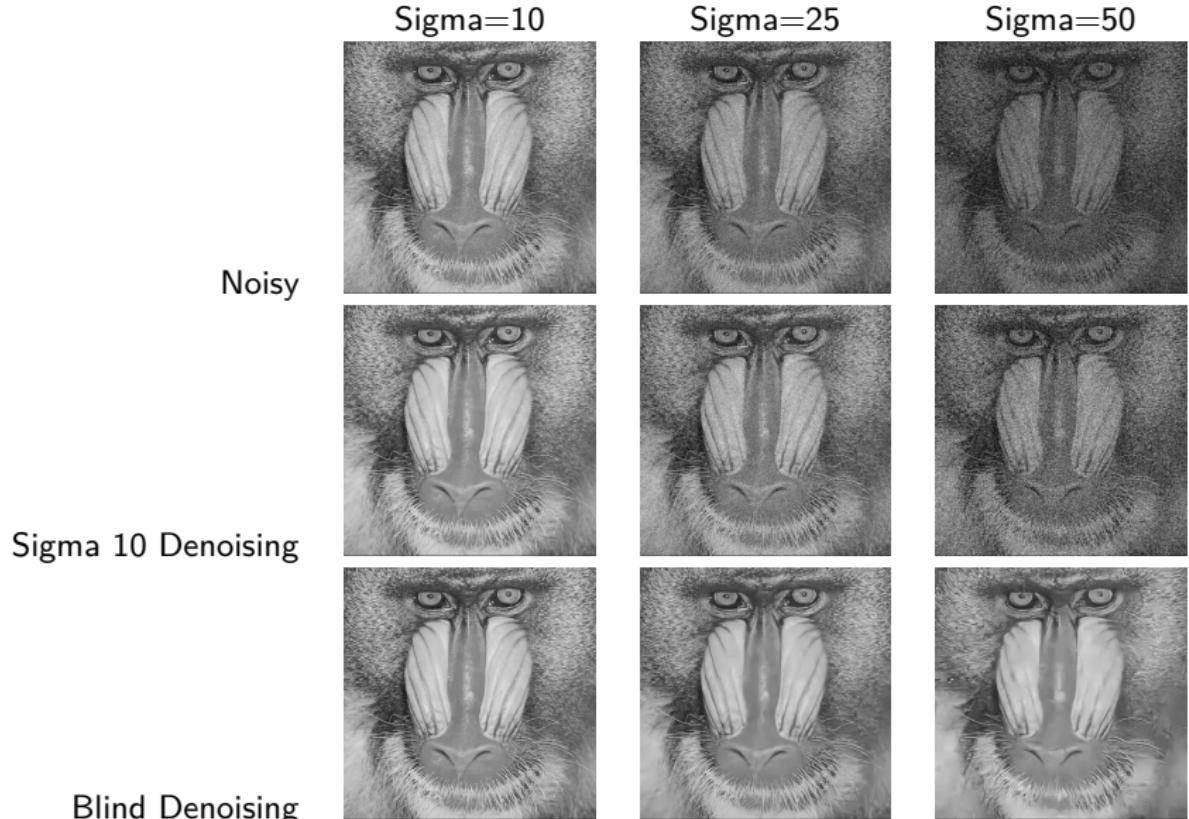
Sigma=30



Sigma=50



Comparison



Comparison

Sigma=10



Sigma=25



Sigma=50



Noisy

Sigma 10 Denoising



Blind Denoising



Comparison

Sigma=10



Sigma=25



Sigma=50



Noisy

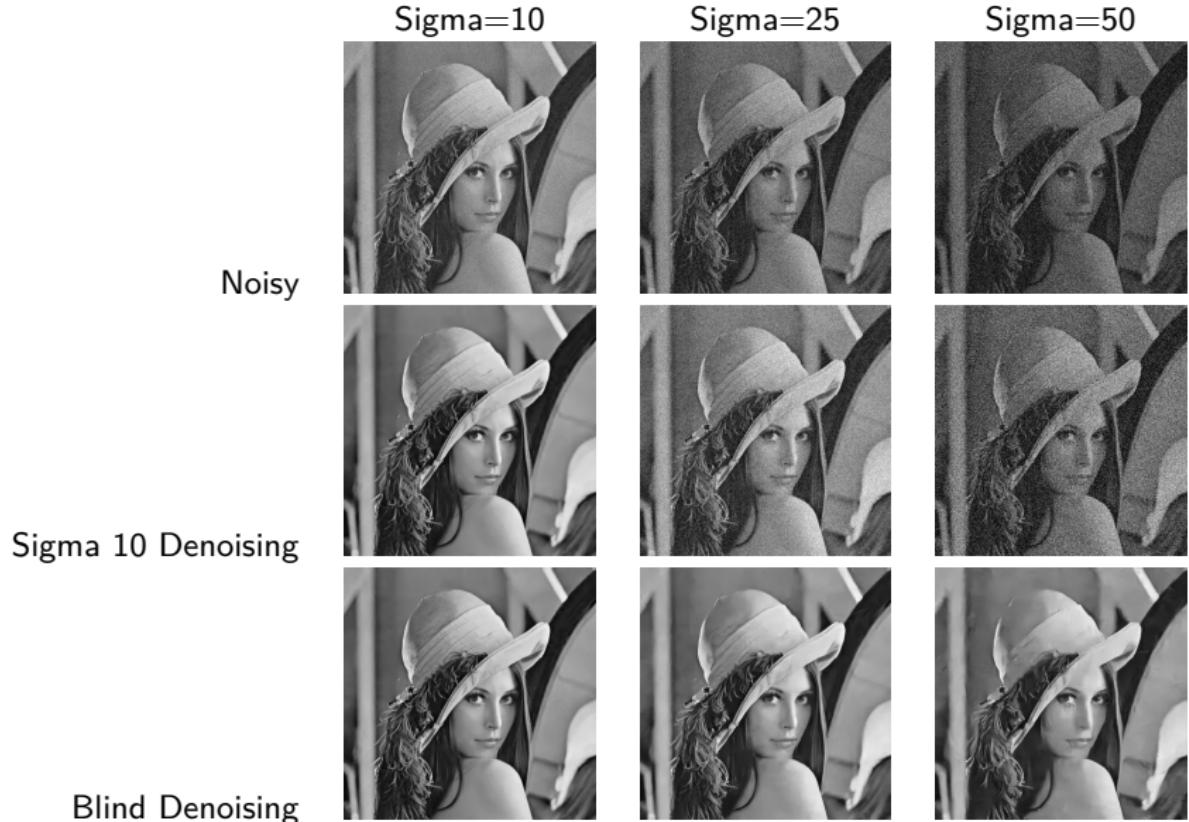
Sigma 10 Denoising



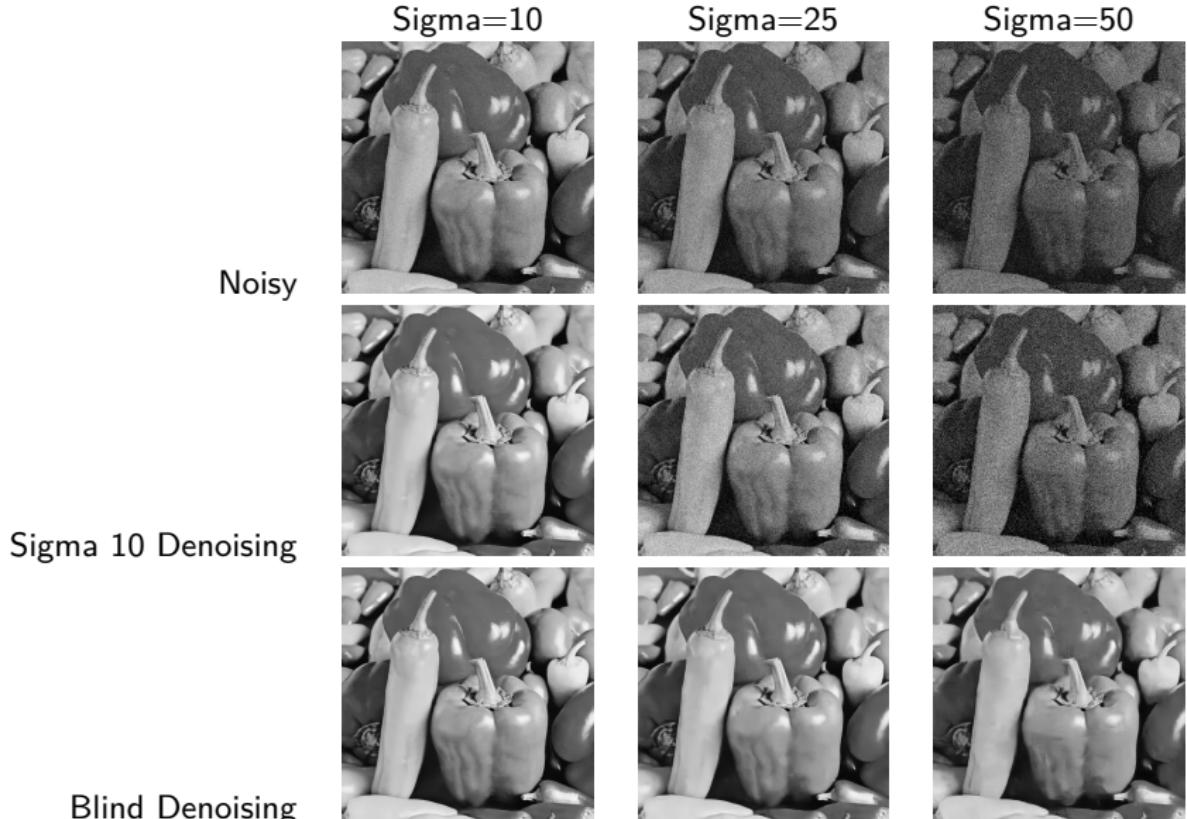
Blind Denoising



Comparison



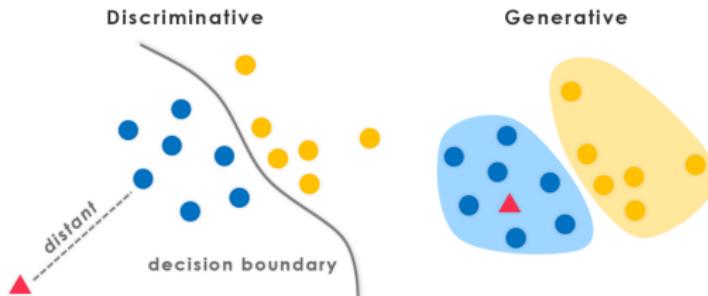
Comparison



References

- 
- Zhang, K., Zuo, W., Chen, Y., Meng, D., & Zhang, L. (2016).
- "Beyond a Gaussian Denoiser: Residual Learning of Deep CNN for Image Denoising."*
- arXiv preprint arXiv:1608.03981.

Generative vs Discriminative



- **Discriminative Approach**
 - Directly finds properties in data.
 - Conditional probabilistic approach $\sim P(\text{class}|\text{data})$
- **Generative Approach**
 - Cares about how the data is generated.
 - Find properties based on the generation.
 - Joint probabilistic approach $\sim P(\text{class}, \text{data})$

Generative Networks



(a) Actual MNIST



(b) Real or Fake?

Outline

1 Denoising

2 Autoencoders

- Traditional Autoencoders
- Regularised Autoencoders

3 Generative Adversarial Networks

4 Improving Discriminator

5 Summary & Interesting Applications

Dimensionality Reduction



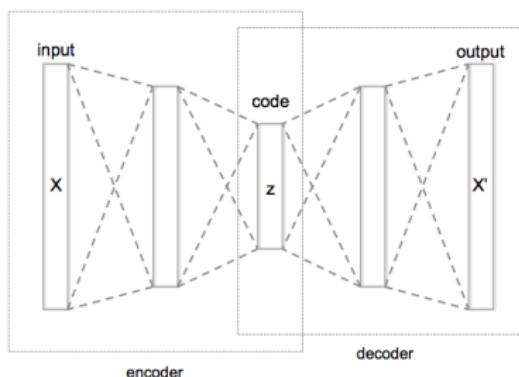
- Extracting emotions from a face.
 - Can reduce to problem of finding relationship with facial features (distance, curvature, etc).
 - Manual selection is difficult.
- PCA is linear \implies limitations in kinds of features extracted
 - top: Original
 - middle: Autoencoder
 - bottom: PCA

Traditional Autoencoders

Aim: $I(X) \approx \psi(\phi(X))$.

- Unsupervised Learning Model

- Let the network learn by itself.
- Encoder: $\phi(X)$
- Decoder: $\psi(z)$
- Learn to generate X by input a smaller dimension z .



Traditional Autoencoders

Aim: $I(x) \approx F(x) = \psi(\phi(x))$.

Cost: $C = C(x, F(x))$

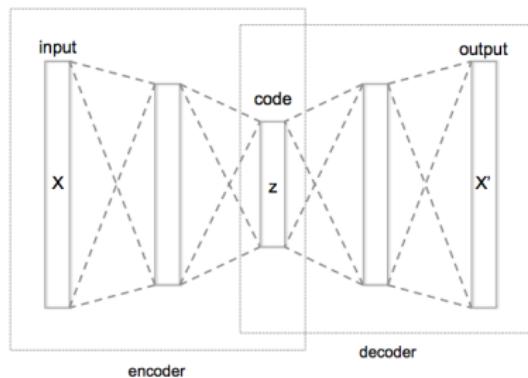
Problem:

With a strong encoding and decoding neural network, $F(x)$ may overfit, without capturing useful properties.

Regularised Autoencoders

Aim: $I(x) \approx F(x) = \psi(\phi(x))$.

Cost: $C = C(x, F(x))$



- Regularised Autoencoder

- Allow model to learn useful features.
- Impose penalties. (Sparse/Contractive autoencoder)
- Augmentation. (Denoising autoencoder)

Outline

1 Denoising

2 Autoencoders

- Traditional Autoencoders
- Regularised Autoencoders
 - Sparse Autoencoder
 - Contractive Autoencoder
 - Denoising Autoencoder

3 Generative Adversarial Networks

4 Improving Discriminator

5 Summary & Interesting Applications

Outline

1 Denoising

2 Autoencoders

- Traditional Autoencoders
- Regularised Autoencoders
 - Sparse Autoencoder
 - Contractive Autoencoder
 - Denoising Autoencoder

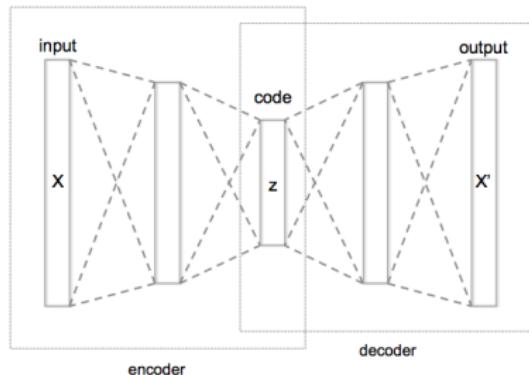
3 Generative Adversarial Networks

4 Improving Discriminator

5 Summary & Interesting Applications

Sparse Autoencoders

$$F(x) = \psi(\phi(x))$$



- Impose penalty to encourage sparseness on code $z = \phi(x)$
- $C(x, F(x) + \lambda \|\phi(x)\|)$

Outline

1 Denoising

2 Autoencoders

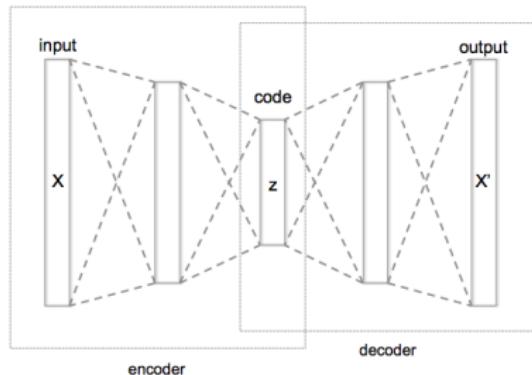
- Traditional Autoencoders
- Regularised Autoencoders
 - Sparse Autoencoder
 - Contractive Autoencoder
 - Denoising Autoencoder

3 Generative Adversarial Networks

4 Improving Discriminator

5 Summary & Interesting Applications

Contractive Autoencoders



- Impose penalty instead, on derivative of z w.r.t. x .
- $C(x, F(x) + \lambda \|\nabla \phi(x)\|)$
- Force model to learn encoder that changes less when x deviates slightly.
 - Locally smooth about the training examples.

Outline

1 Denoising

2 Autoencoders

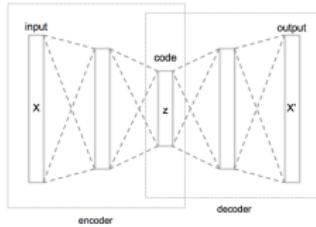
- Traditional Autoencoders
- Regularised Autoencoders
 - Sparse Autoencoder
 - Contractive Autoencoder
 - Denoising Autoencoder

3 Generative Adversarial Networks

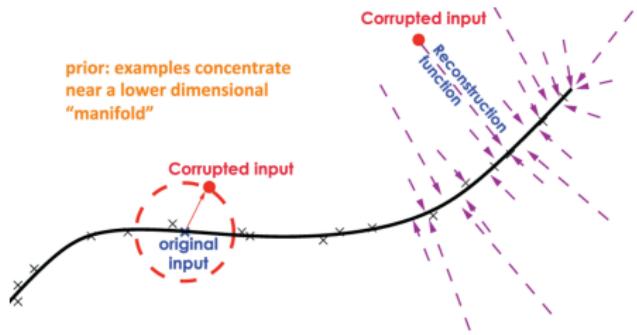
4 Improving Discriminator

5 Summary & Interesting Applications

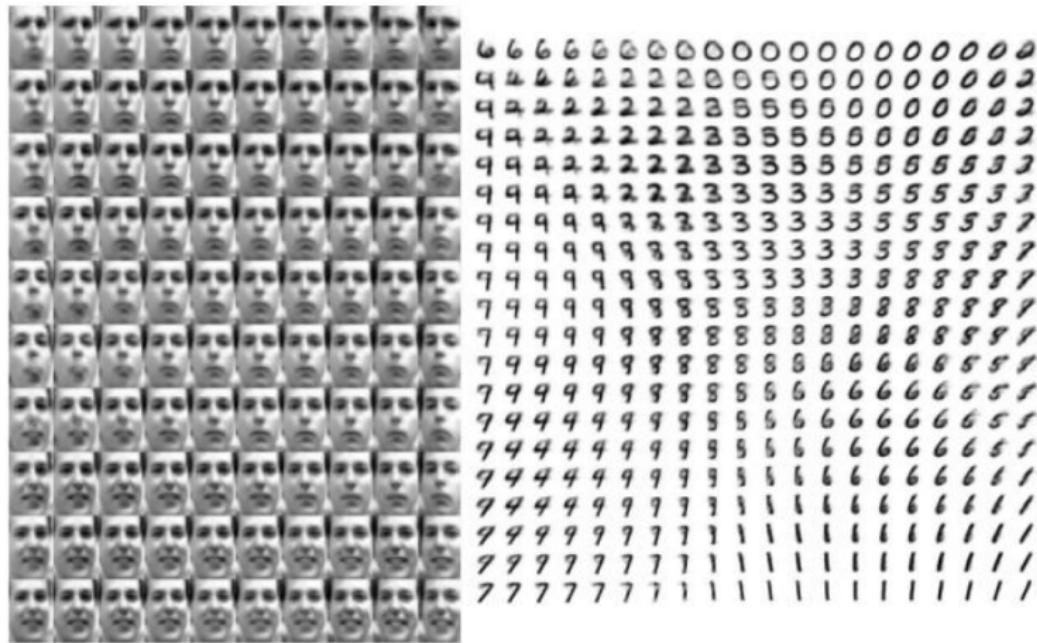
Denoising Autoencoders



- Augment input x by some noise ϵ .
- $C(x, F(x + \epsilon))$



2D Manifold of MNIST Digits



2D-Manifold
y-axis: Emotions Digits
x-axis: Rotation

References

-  Buduma, N. "*The curse of dimensionality and the autoencoder.*" <http://nikhilbuduma.com/2015/03/10/the-curse-of-dimensionality>
-  Bengio, Y., Goodfellow, I. J., & Courville, A. (2015). "*Deep learning.*"
-  Pascal Vincent, Hugo Larochelle, Yoshua Bengio, and Pierre-Antoine Manzagol. "*Extracting and composing robust features with denoising autoencoders.*" In ICML, 2008.
-  Diederik P Kingma and Max Welling. "*Auto-Encoding Variational Bayes*" ICLR, 2014.
-  Doersch, C. (2016). "*Tutorial on variational autoencoders.*" arXiv preprint arXiv:1606.05908.

Outline

1 Denoising

2 Autoencoders

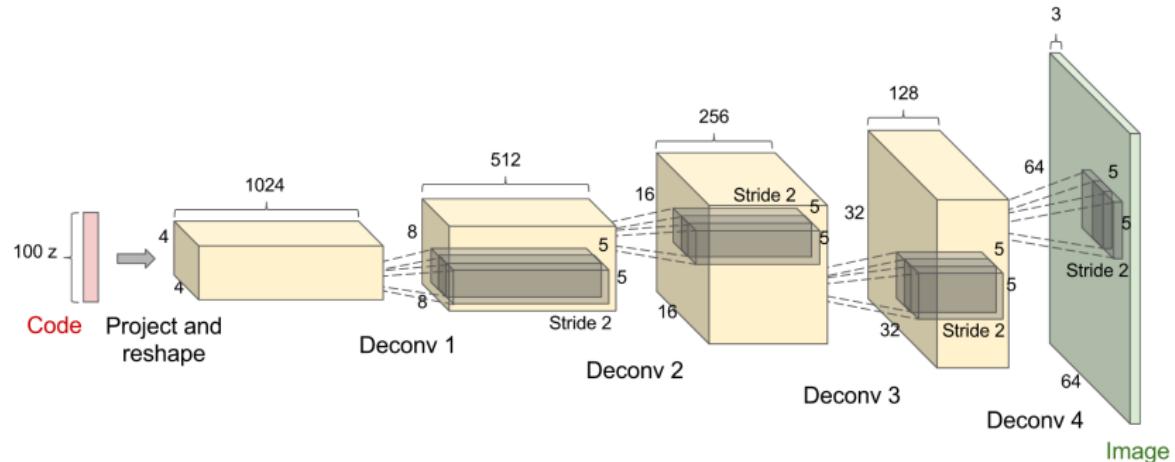
3 Generative Adversarial Networks

- Generative Adversarial Network (GAN)
- Experiment on MNIST
- Transposed Convolution

4 Improving Discriminator

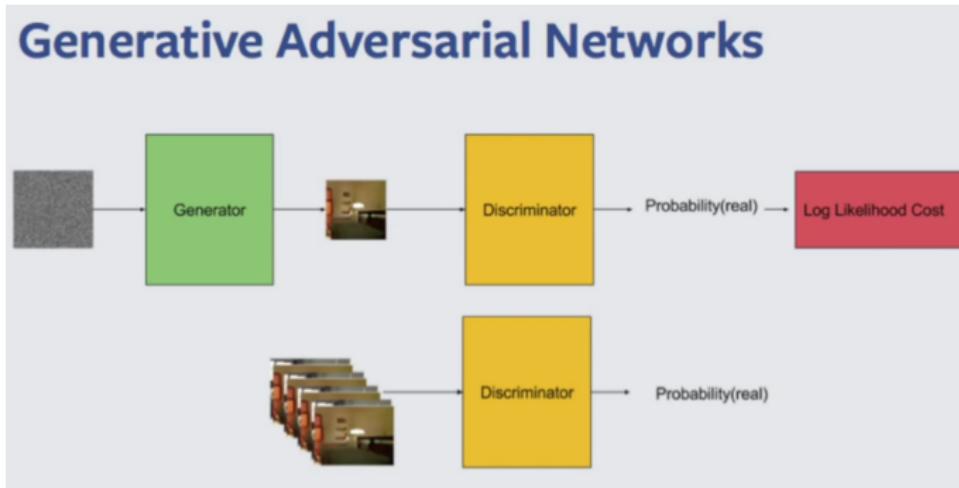
5 Summary & Interesting Applications

Constructing Loss



- Want generated images to look real.
- How to construct loss function?

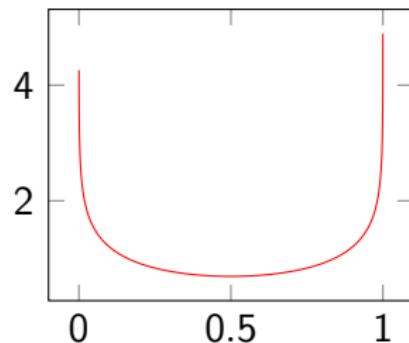
Generative Adversarial Network (GAN)



Estimating generative models via an adversarial process

- Generative model G
 - Aims to have D classify wrongly.
 - **Learns data distribution.**
- Discriminative model D
 - Learn input is real or fake data → binary problem.
 - **Unable to distinguish real or fake at the end.**

Cross Entropy Cost



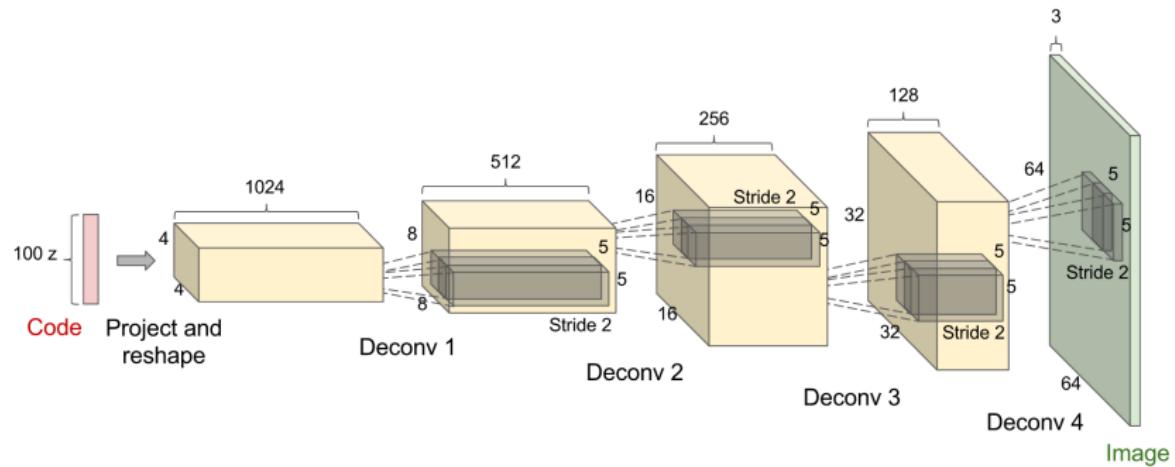
$$C(\mathbf{x}, \mathbf{w}, \mathbf{b}) = - \left[\left(y \log a^L \right) + (1 - y) \log (1 - a^L) \right]$$

Generative Adversarial Network (GAN)

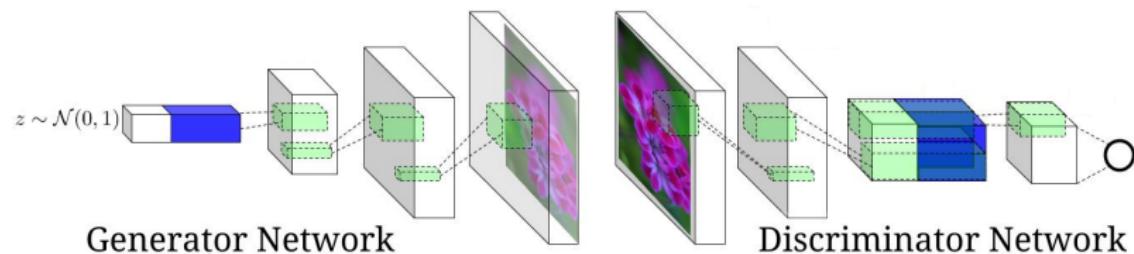
$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim P_X(x)} [\log D(x)] + \mathbb{E}_{z \sim P_Z(z)} [\log(1 - D(G(z)))]$$

- $D(x)$
 - $D(x) \approx 1 \implies x$ is likely to be data.
 - $D(x) \approx 0 \implies x$ is likely to be artificial.
- $G(z)$
 - With code z , generates image $G(z)$ that attempts to look like X .
- $V(D, G)$
 - If quadratic:
$$V(D, G) = \mathbb{E}_{x \sim P_X(x)} \left[-\|D(x) - 1\|_2^2 \right] + \mathbb{E}_{z \sim P_Z(z)} \left[-\|D(G(z))\|_2^2 \right]$$
 - Experimentally, output looks blurry

Generator



Generator & Discriminator



- Hybrid of Dense and Convolutional layers

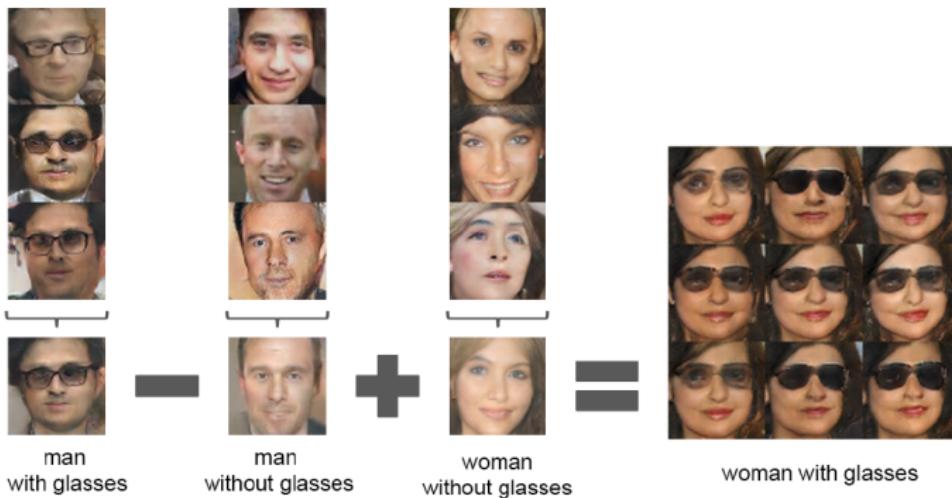
Generative Adversarial Network (GAN)

Autoencoder vs GAN

- Autoencoder: Aim to get pixel-perfect reconstruction
- GAN: Aim to construct believable images that fit the distribution

Generative Adversarial Network (DCGAN)

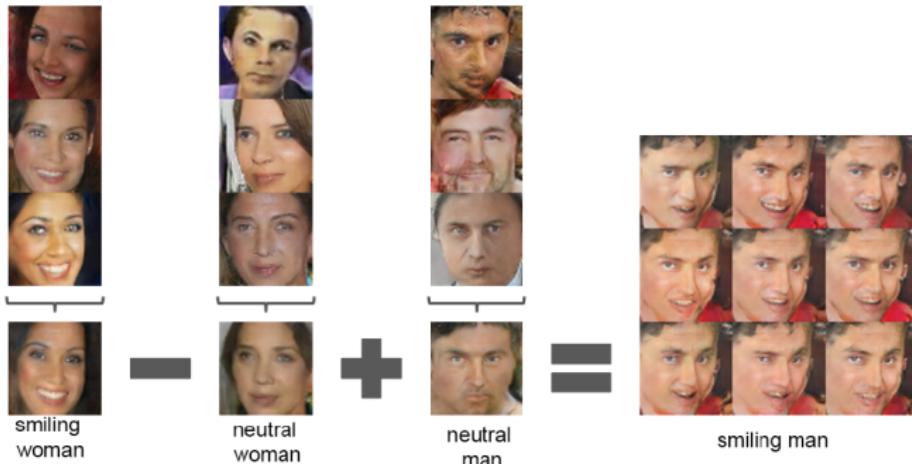
Vector arithmetic for visual concepts.



Locally smooth generator property -
perturb mean by uniform noise $\sim U(-0.25, 0.25)$

Generative Adversarial Network (DCGAN)

Vector arithmetic for visual concepts.



Locally smooth generator property -
perturb mean by uniform noise $\sim U(-0.25, 0.25)$

Generative Adversarial Network (DCGAN)

Interpolation.



Figure 8: A "turn" vector was created from four averaged samples of faces looking left vs looking right. By adding interpolations along this axis to random samples we were able to reliably transform their pose.

References

- [GAN] Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., ... & Bengio, Y. (2014). “*Generative adversarial nets.*” In Advances in Neural Information Processing Systems (pp. 2672-2680).
- [LAPGAN] Denton, Emily L., Soumith Chintala, and Rob Fergus. “*Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks.*” Advances in neural information processing systems. 2015.
- [DCGAN] Radford, A., Metz, L., & Chintala, S. (2015). “*Unsupervised representation learning with deep convolutional generative adversarial networks.*” arXiv preprint arXiv:1511.06434.
- “From Facebook AI research - Soumith Chintala - Adversarial Networks”
(Youtube Link)

Outline

1 Denoising

2 Autoencoders

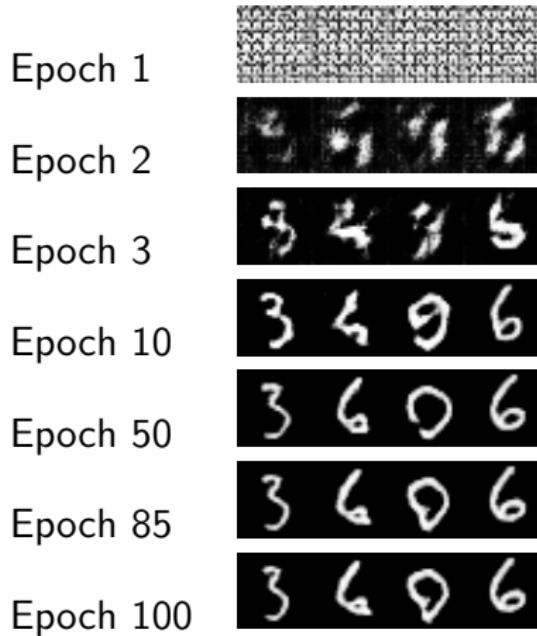
3 Generative Adversarial Networks

- Generative Adversarial Network (GAN)
- Experiment on MNIST
- Transposed Convolution

4 Improving Discriminator

5 Summary & Interesting Applications

Experiment on MNIST



Outline

① Denoising

② Autoencoders

③ Generative Adversarial Networks

- Generative Adversarial Network (GAN)
- Experiment on MNIST
- Transposed Convolution
 - Padding
 - Strides

④ Improving Discriminator

⑤ Summary & Interesting Applications

Outline

1 Denoising

2 Autoencoders

3 Generative Adversarial Networks

- Generative Adversarial Network (GAN)
- Experiment on MNIST
- Transposed Convolution
 - Padding
 - Strides

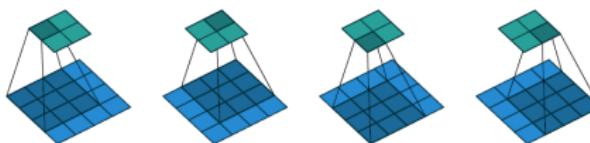
4 Improving Discriminator

5 Summary & Interesting Applications

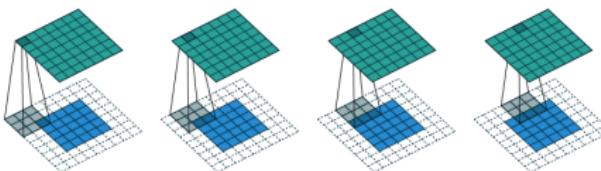
Typical Convolution: Padding

For a $m \times m$ filter f , consider $A * f$

- Valid
 - No padding.



- Full
 - Pads A by $(m-1) \times (m-1)$.

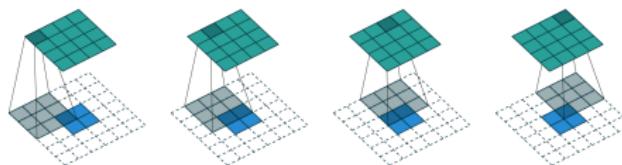


- Same
 - Pads A by $\frac{m-1}{2} \times \frac{m-1}{2}$ to get same output size.

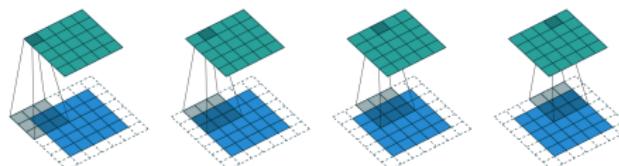
Transposed Convolution: Padding

Visualisation: Equivalent to its “dual” original convolution

- ‘Valid’ Transposed Conv on 2×2 input
 - Use ‘Full’ with $*$ operation



- ‘Same’ Transposed Conv with 3×3 filter
 - Use ‘Same’ with $*$ operation



Outline

1 Denoising

2 Autoencoders

3 Generative Adversarial Networks

- Generative Adversarial Network (GAN)
- Experiment on MNIST
- Transposed Convolution
 - Padding
 - Strides

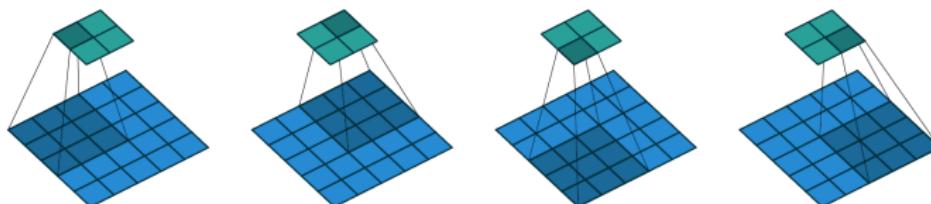
4 Improving Discriminator

5 Summary & Interesting Applications

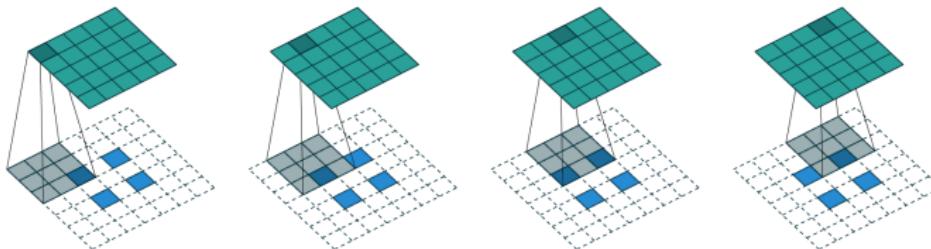
Transposed Convolution: Strides

Insert zeros in between pixels (Dilation)

- Strides $s: 2 \times 2$



- Transposed Conv Strides $s': 1 \times 1$
- Insert: 1×1



Summary

Convolution described by

- filter size: m
- stride: s
- padding: p

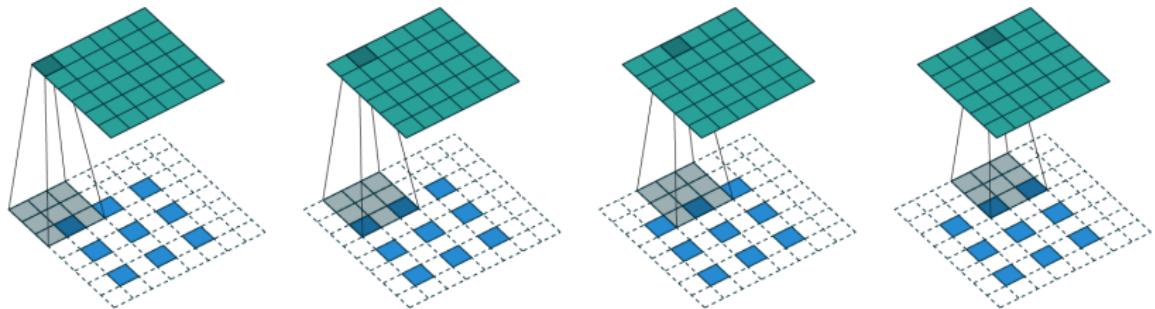
has associated Transposed Convolution using

- filter size: $m' = m$
- stride: $s' = 1$
- padding: $p' = m - p - 1$
- insert/dilation: $i' = s - 1$
- diag: $a' = (i + 2p - m) \bmod s$

with output size: $o' = s'(S - 1) + a' + m' - 2p'$,

where the 'stretched size' S is after adding i' zeros between each pixel

Summary: Example



Backward

Input: 6×6

Filter: 3×3

Stride: 2×2

Padding: 1×1

Transposed Conv

Input: 3×3

Filter: 3×3

Stride: 1×1

Padding: 1×1

Insert: 1×1

Diag: 1×1

Visualisation

References



Dumoulin, V. and Visin, F., 2016. *"A guide to convolution arithmetic for deep learning."* arXiv preprint arXiv:1603.07285.

Outline

- 1 Denoising
- 2 Autoencoders
- 3 Generative Adversarial Networks
- 4 Improving Discriminator
 - Adversarial Data
- 5 Summary & Interesting Applications

Fooling ConvNets

Error in classification with minor perturbation.



(a) Correct Classification (b) Perturbation (c) Wrong Classification

- Perturb image by a small magnitude of $\frac{\partial C(x; \theta)}{\partial x}$
 - Gradient **Ascent**: $\tilde{x} = x + \epsilon \frac{\partial C(x; \theta)}{\partial x}$
- Note: Up until now, we have been discussing $\frac{\partial C(x; \theta)}{\partial \theta}$ instead
 - Gradient **Descent**: $\theta_{k+1} = \theta_k - \eta \frac{\partial C(x; \theta_k)}{\partial \theta_k}$

Improving Discriminator

- Impose smooth local behaviour
 - Locally constant prior instead to perturb.
- Discriminator is more robust to perturbations
 - Assume different classes lies in disconnected manifolds.
 - Aim: $F(\hat{x}) = F(\hat{x} + \epsilon)$, for unlabelled \hat{x} to lie locally in the same class.
 - Perturbations should not jump to other manifolds easily.
 - Better distinction of classes.
- To control $\|\tilde{x} - x\|_\infty$
 - Use $\text{sign}\left(\frac{\partial C(x; \theta)}{\partial x}\right)$ instead to perturb.
- Redefine Cost function: $\hat{C}(x; \theta) = \alpha C(x; \theta) + (1 - \alpha) C(\tilde{x}; \theta)$
- A form of data augmentation: increase training data set adaptively.

References

-  Goodfellow IJ, Shlens J, Szegedy C. "*Explaining and harnessing adversarial examples.*" arXiv preprint arXiv:1412.6572. 2014 Dec 20.
-  Bengio, Y., Goodfellow, I.J. & Courville, A., 2015. "*Deep learning.*" An MIT Press book in preparation.

Outline

- 1 Denoising
- 2 Autoencoders
- 3 Generative Adversarial Networks
- 4 Improving Discriminator
- 5 Summary & Interesting Applications
 - Recent Progress

Summary

- Neural Networks (Compositionality)
- Gradient Descent (Variants)
- Backpropagation
- Initialisation (Weights, Hyperparameters)
- Regularisation (Cost, Layers)
- Convolutional Neural Networks (Design, Architectures)
- Discriminators & Generators

Some Recent Progress

- Images
 - Image super-resolution
 - Object detection and classification
 - Detection and text-recognition
- Videos
 - Sharpening live videos
 - Fake Video
 - Modifying Trump's Hair from a video (26:00)
- Audio & Videos
 - Creating realistic-audio effects from silent video

Thank You